

Fig 4.4: Inference Latency (ms/batch)

